

1. Write a program to implement CNN.

CNN-MNIST

ANSWER # Step 1: Import necessary libraries

```
import tensorflow as tf
```

```
from tensorflow.keras import layers, models
```

```
from tensorflow.keras.datasets import mnist
```

```
import matplotlib.pyplot as plt
```

Step 2: Load and preprocess the MNIST dataset

Load the dataset from keras

```
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

Reshape the data to add a color channel dimension

```
train_images = train_images.reshape(60000, 28, -28, 1)
```

```
test_images = test_images.reshape(10000, 28, 28, 1)
```

Normalize the pixel values to range from 0 to 1

```
train_images, test_images = train_images / 255.0, -test_images / 255.0
```

Step 3: Build the CNN model

```
model = models.Sequential()
```

Add a convolutional layer with 32 filters, 3x3-kernel size, and ReLU activation

```
model.add(layers.Conv2D(32, (3, 3), activation = -'relu', input_shape = (28, 28, 1)))
```

Add a max-pooling layer

model.add(layers.MaxPooling2D((2,2)))

Add a second convolutional layer

model.add(layers.Conv2D(64, 3, 3, activation=

- = 'relu'))

Add a max-pooling layer

model.add(layers.MaxPooling2D((2,2)))

Add a third layer convolutional layer

model.add(layers.Conv2D(64, 3, 3, activation='relu'))

Flatten the output for the dense layers

model.add(layers.Flatten())

Add a fully connected layer (Dense layer)

model.add(layers.Dense(64, activation='relu'))

Add the output layer with 10 neurons (one-

-for each digit) and softmax activation.

model.add(layers.Dense(10, activation=~~'softmax'~~
- 'softmax'))

Step 4: compile the model

model.compile(optimizer='adam', loss='sparse-
categorical_crossentropy', metrics=['accuracy'])

Step 5: Train the model

history = model.fit(train_images, train_labels, epochs=

$-=5$, batch_size=64, validation_split=0.1)

Step 6: Evaluate the model

```
test_loss, test_acc = model.evaluate(test_images, -  
-test_labels)
```

print

print(f'Test accuracy: {test_acc}')

Step 7: Visualize the accuracy and loss over epochs

```
plt.plot(history.history['accuracy'], label='accuracy')  
plt.plot(history.history['val_accuracy'], label='val-a-  
ccuracy')
```

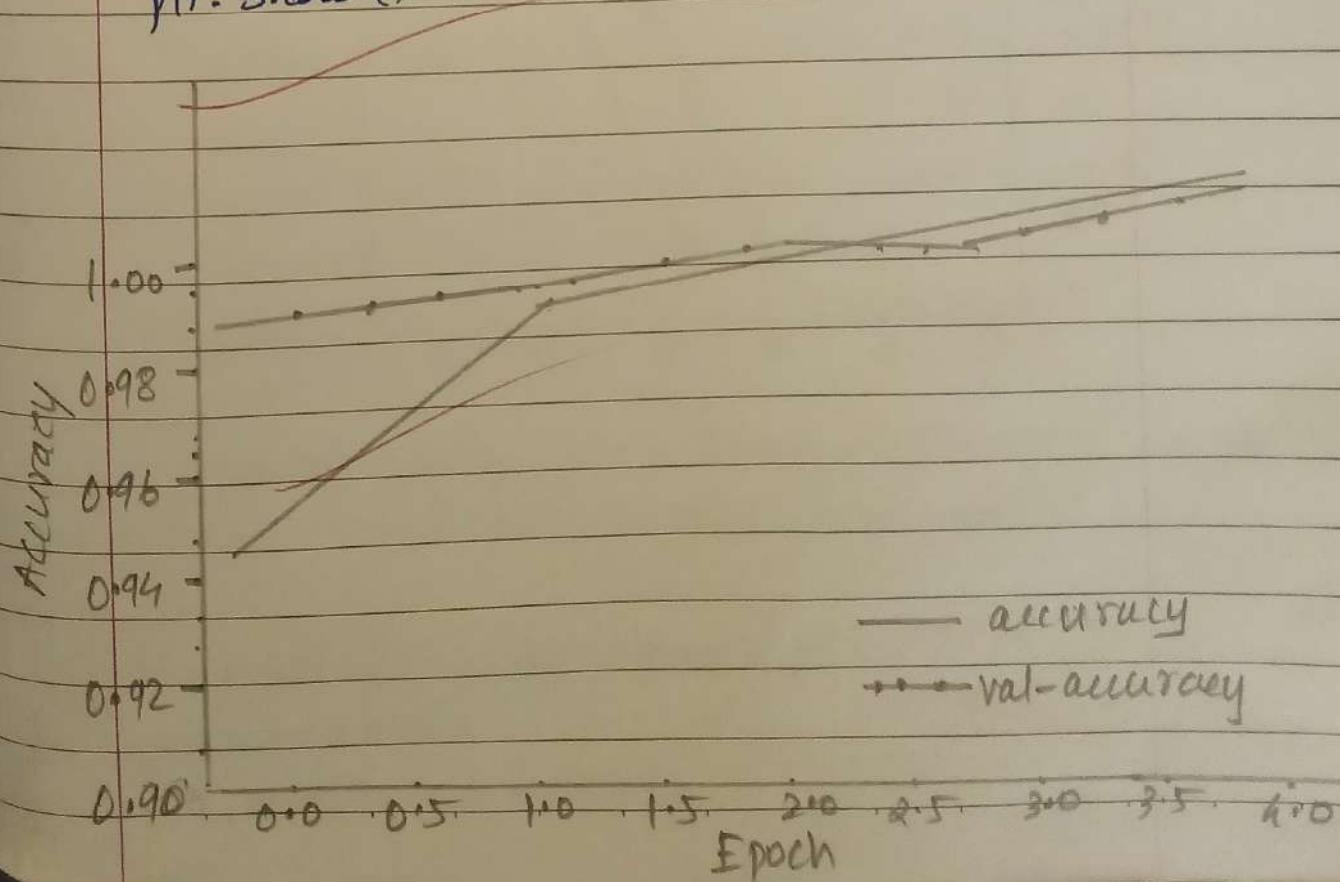
plt.xlabel('Epoch')

plt.ylabel('Accuracy')

plt.ylim([0.9, 1])

plt.legend(loc='lower right')

plt.show()



```
# Step 8: Make predictions on test data  
predictions = model.predict(test_images)
```

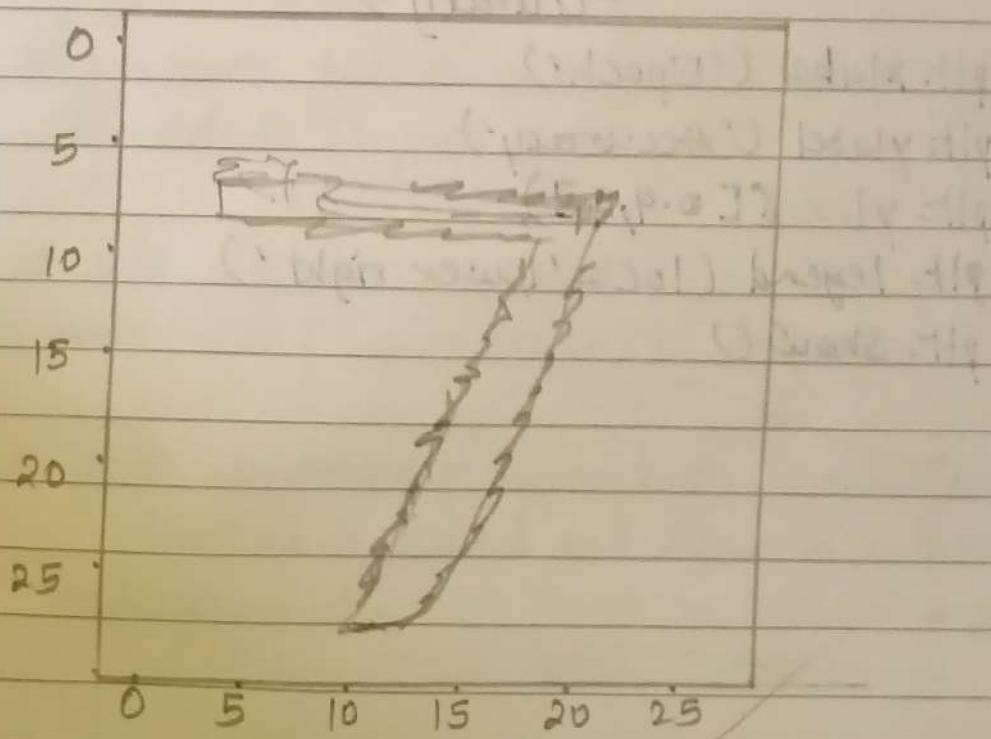
```
# Visualize the first test image and its predicted-  
-label
```

```
plt.imshow(test_images[0].reshape(28,28),  
           cmap=plt.cm.binary)
```

```
plt.title(f'Predicted label = {predictions[0].argmax()}\naccuracy = {accuracy}')
```

plt.show()

Predicted Label: 7



2. COPY-of-DL2

Ans-

```

import pandas as pd
from keras.layers import Dense
from keras.models import Sequential
from google.colab import files
uploaded = files.upload()

dataset = pd.read_csv('diabetes.csv')
dataset.describe()

X = dataset.drop(['Outcome'], axis=1)
y = dataset['Outcome']

model = Sequential()
model.add(Dense(12, input_dim=8, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

model.fit(X, y, epochs=150, batch_size=10)

accuracy = model.evaluate(X, y)
print("Model accuracy is ", accuracy)

model.compile(loss='binary_crossentropy', optimizer
              = 'adam', metrics=['accuracy'])

prediction = model.predict(X)
model.summary()

```

3. Data Analytics

AUSY # Sample Machine learning Project
 # Dataset: grapes-new.csv

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import statsmodels.api as sm
import seaborn as sns
```

```
from google.colab import files
uploaded = files.upload()
```

```
df = pd.read_csv('grapes-new.csv')
```

Dataset is now stored in a pandas DataFrame

df

df.columns

Separate input-output data

x = df.drop('CLASS', axis=1)

y = df['CLASS']

x.shape

x.info

Multi class classification

y.value_counts()

2. Data cleaning-Missing values

x.isnull().sum()

x.fillna(method='backfill', inplace=True)

x.fillna(method='pad', inplace=True)

x.isnull().sum()

3. Exploratory data analysis

x.describe()

Label Encode

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
```

```
x['SHAPE'] = le.fit_transform(x['SHAPE'])
```

```
x['SOIL-TYPE'] = le.fit_transform(x['SOIL-TYPE'])
```

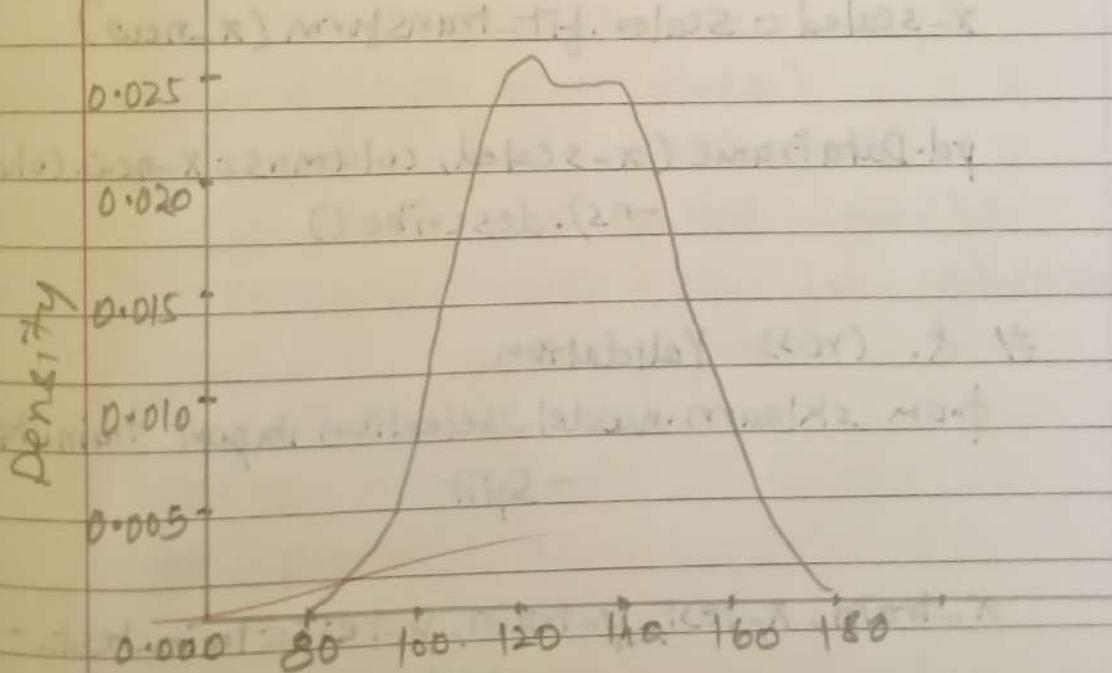
```
(x['SHAPE'])
```

```
x['SOIL-TYPE'] = le.fit
```

```
x['SOIL-TYPE'] = le.fit_transform(x['SOIL-TYPE'])
```

```
x.describe()
```

```
sns.kdeplot(df['COLOR-INTENSITY'])
```



COLOR-INTENSITY

```
temp = df[['COLOR-INTENSITY', 'RIPENESS-PER', -  
          'ALCOHOL-PER', 'FLAVANOIDS', 'FERT-NITRO-PER',  
          'CLASS']]
```

```
sns.pairplot(temp, hue='CLASS', palette='tab10')
```

6. Feature selection.

```
from sklearn.feature_selection import SelectKBest, chi2
```

```
skf = SelectKBest(score_func=chi2, k=5)
```

```
x_new = skf.fit_transform(x, y)
```

$x_{\text{new}}. \text{shape}$

```
skf.get_support()
```

```
x_new = x.iloc[:, skf.get_support()]
```

x_{new}

6. Feature scaling

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
```

```
x_scaled = scaler.fit_transform(x_new)
```

```
pd.DataFrame(x_scaled, columns=x_new.columns).describe()
```

8. Cross Validation

```
from sklearn.model_selection import train_test_split
```

$x_{\text{train}}, x_{\text{test}}, y_{\text{train}}, y_{\text{test}} = \text{train_test_split}(x_{\text{scaled}}, y, \text{random_state}=0)$

$x_{\text{train}}. \text{shape}$

$x_{\text{test}}. \text{shape}$

9. Algorithms, Hyperparameter Tuning

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.model_selection import GridSearchCV

```

param = {

- 'random_state': [0, 1, 2, 3, 4, 5],
- 'criterion': ['gini', 'entropy'],
- 'splitter': ['best', 'random']

}

dt = DecisionTreeClassifier()

grid = GridSearchCV(dt, param_grid=param, cv=5, scoring='accuracy')

grid.fit(x_train, y_train)

y_pred = grid.predict(x_test)

accuracy_score(y_test, y_pred)

ConfusionMatrixDisplay.from_predictions(y-test, y-pred)

12. Deployment

.Serialization

import joblib

4. Data Analysis Using Pandas

Ans → # Find your working directory

```
import os
os.getcwd()
os.getcwdu()
```

```
import pandas as pd
df = pd.read_csv('student.csv')
```

df

type(df)

If no headers are present then first row will be considered as header

```
df1 = pd.read_csv('student1.csv')
df1
```

If we want to add headers then add names like below

```
df1 = pd.read_csv('student1.csv', names=-
                  ['Rno', 'Name', 'Marks', 'Age'])
```

df1

If data is not separated we will have to add separator

```
df3 = pd.read_csv('student3.csv')
df3
```

Separator added

`df3 = pd.read_csv('student3.csv', sep='|')`

`df.shape`

`df.ndim`

`df.columns`

`df.dtypes`

`df.T # transpose`

Display selected rows

`df.head(3)`

Display selected rows

`df.tail(3)`

iloc index location

`df.iloc[2:8, 1:4]`

`df.iloc[2:, :4] # before comma row, after -
- comma column`

~~`df.iloc[2:8, :4] # before comma row`~~

~~`df.iloc[:, :] # all the rows & all the columns`~~

`x = [1, 2, 3]`

`y = x`

~~`x.append(4)`~~

~~`x`~~

~~`y`~~

~~`df.iloc[3, :]`~~

`x = df.iloc[:, 1]`

`type(x)`

`df.iloc[3, 1]`

`df.iloc[[3, 6, 8, 1], :] # selected rows by all columns`

df.iloc[3,:]

df.iloc[[3,6,8,1],[1,4]] # Selected rows - by selected columns

df.loc[:,['name','age']]

df.name # one column all rows

df[['name','marks']] # selected columns - all rows

df.iloc[[3],:] # data frame

df.iloc[3,:] # LIST

df3.drop(5)

df.drop([5,7,2])

df.drop(range(0,10,2)) # Display odd no. rows

df.drop('class',axis=1) # Drop column

df.drop(['class','age'],axis=1) # Axis 1 - all rows drop to work on column

df['roll'] + 100

df['marks'] = df['marks'] + 0.5

df

df['cgpa'] = df['marks'] / 10 + 0.38

x = ['M','M','M','M','F','F','M','F','M','M']

df['gender'] = x

df

df = pd.read_csv('student.csv')

df

ndf = pd.DataFrame({})

'roll': range(1,11),

'cgpa': df['marks'] / 10 + 0.38

'gender': x

})

ndf

df.columns

ndf.columns

df.merge(ndf, on='roll') # Merges two data

df5 = pd.read_csv('student5.csv')

df5

pd.concat([df, df5])

pd.concat([df, df5], ignore_index=True)

~~df~~

df['name'].str.upper()

df['name'].str.title()

df['name'].str.lower()

df['name'].str.replace('id', 'ee')

df.info()

df.describe()

df['class'].mode()

df['age'].mode()

~~df.sum()~~

df.sum(axis=1)

df['marks'] > 60

df[df['marks'] > 60]

x = df[(df['marks'] > 60) & (df['class'] == 'BE')]

df[df['name'].str.startswith('an')]

Data Export

x.to_csv('output.csv', index=False)

x.to_html('output.html', index=False)

5. Data Preprocessing

Ans → Data Cleaning :-

import pandas as pd

df = pd.read_csv('student3.csv')

df

df['age'].mean()

df['marks'].mean()

fix structural error... remove spaces before
and after data... trim spaces in column name
strip is applied only on string data

df.columns = df.columns.str.strip()

df.columns

df

applied on categorical data

df['class'].value_counts()

df['class'] = df['class'].str.upper()

df['class'].value_counts()

df['age'].value_counts()

df['name'] = df['name'].str.title()

df

checks duplicate entries

df.duplicated()

df[df.duplicated()]

df.duplicated().sum()

df.drop_duplicates(inplace=True)

df

Handling missing values

df.info()

df.count()

df.isnull()

Find the count of df missing values

df.isnull().sum()

df.describe()

1. drop all the rows with missing values

df.dropna()

2. Missing values as per our handling knowledge,
fix one value for column and replace all missing
values

df['age'].fillna(21)

3. Replace missing values Mean - for continuous

df['marks'].fillna(df['marks'].mean())

3. Median mode

df['class'].fillna(df['class'].mode()[0])

df.fillna(method='pad')

df.fillna(method='backfill')

Data Encoding :

Convert string data to number

Label encoding

df = pd.read_csv('student.csv')

df

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
```

```
df['name'] = le.fit_transform(df['name'])
df
```

```
# One Hot Encoding
pd.get_dummies(df)
```

Data Scaling / Feature scaling

```
df = pd.read_csv('social-Network-Ads.csv')
df
```

```
x = df[['Age', 'EstimatedSalary']]
```

```
import matplotlib.pyplot as plt
plt.scatter(x['Age'], x['EstimatedSalary'], c=df['Purchased'])
```

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
```

```
x_scaled = scaler.fit_transform(x)
```

```
x_scaled = pd.DataFrame(x_scaled, columns=x.columns)
```

x_scaled

x_scaled.describe()

x.std()

```
from sklearn.preprocessing import StandardScaler
```

M	T	W	T	F	S	S
Page No.:	YOUVA					
Date:						

- Scalar

scalar = StandardScaler()

x-scaled = scalar.fit_transform(x)

x-scaled = pd.DataFrame(x-scaled, columns=x.
columns)

x-scaled.std()

plt.scatter(x['Age'], x-scaled['EstimatedSalary'],
c=df['Purchased'])

6. Decision Tree Algorithms

```

Ans → # Import necessary libraries
from sklearn.datasets import load_iris
from sklearn.model_selection import train-
    -test-split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
from sklearn import tree
import matplotlib.pyplot as plt

```

```

# Load the Iris dataset
iris = load_iris()
X = iris.data # Features
y = iris.target # Target Labels

```

Split the dataset into training and testing sets
~~X-train, X-test, y-train, y-test = train-test-split -
 - (X, y, test_size=0.3, random_state=42)~~

Initialise the Decision Tree classifier
~~clf = DecisionTreeClassifier(max_depth=3, rand-~~
~~-om_state=42)~~

Train the model
~~clf.fit(X-train, y-train)~~

Predict on the test set
~~y-pred = clf.predict(X-test)~~

Evaluate the model's accuracy

accuracy = accuracy_score(y-test, y-pred)

print(f"Accuracy: {accuracy * 100:.2f}%")

Visualize the decision tree

plt.figure(figsize=(12,8))

tree.plot_tree(clf, feature_names=iris.feature_names, class_names=iris.target_names, filled=True)

plt.show()

Display the classification report

from sklearn.metrics import classification_report

print(classification_report(y-test, y-pred, target_names=iris.target_names))

7. Fashion CNN

ANSWER

```
# importing the libraries
```

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn.model_selection import train_
    -test-split
```

```
from sklearn.metrics import accuracy_score
```

```
from tqdm import tqdm
```

```
import torch
```

```
from torch.autograd import Variable
```

```
# newly added for FashionMNIST Dataset
```

```
from torch.utils.data import Dataset
```

```
from torchvision import datasets
```

```
from torch.nn import Linear, ReLU, CrossEnt-
```

```
-ropyLoss, Sequential, Conv2d, MaxPool2d, Module,
```

```
Softmax, BatchNorm2d, Dropout
```

```
from torch.optim import Adam, SGD
```

```
import torch
```

```
import torch.nn as nn
```

```
import torchvision
```

```
import torch.nn.functional as F
```

```
import torchvision.transforms as transforms
```

```
import matplotlib.pyplot as plt
```

''' Check whether a CUDA enabled GPU is

available else use CPU ""

device = torch.device('cuda' if torch.cuda.is_

-available() else 'cpu')

transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5), (0.5))])

" " Using the EMNIST database has broken the program using FashionMNIST dataset instead " "

training_data = datasets.FashionMNIST(root=' - contents /', download=True, transform=trans-
form, train=True)

test_data = datasets.FashionMNIST(root='cont-
ents /', download=True, transform=transform,
train=False)

train_loader = torch.utils.data.DataLoader(data-
set=training_data,
batch_size=128,
shuffle=True)

test_loader = torch.utils.data.DataLoader(data-
set=test_data,
batch_size=128,
shuffle=True)

for i in range(6):

plt.subplot(2, 3, i+1)

plt.imshow(training_data.data[i])

plt.title('Ground Truth: {}'.format(training-
data.classes[training_data.targets[i]]))

plt. axis('off')

class CNNModel(nn.Module):

def __init__(self):

super(CNNModel, self).__init__()

self.conv1 = nn.Conv2d(in_channels=1,
out_channels=32, kernel_size=5, stride=1)

self.conv2 = nn.Conv2d(32, 64, 5, 1)

self.fc = nn.Linear(64 * 20 * 20, 47)

def forward(self, x):

x = F.relu(self.conv1(x))

x = F.relu(self.conv2(x))

8. K-Means Clustering

ANSWER ~~Not~~ Import Libraries

Import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

df = pd.read_csv('Mall-Customers.csv')

df

x = df.iloc[:, 3:]

x

plt.title('Data Distribution')

plt.xlabel('Annual Income')

plt.ylabel('Spending Score')

plt.scatter(x['Annual Income (k\$)'], x['Spending Score (1-100)'])

from sklearn.cluster import KMeans

km = KMeans(n_clusters=3)

labels = km.fit_predict(x)

labels

km.inertia_

~~Elbow Method~~

sse = []

for k in range(1, 16):

km = KMeans(n_clusters=k)

km.fit_predict(x)

sse.append(km.inertia_)

sse

```
plt.plot(sse, color='red', marker='o', mfc='blue')
```

```
from kneed import KneeLocator
k1 = KneeLocator(range(1, 16), sse, curve='convex', direction='decreasing')
k1.elbow
```

Apply K-means cluster

```
km = KMeans(n_clusters=5)
```

```
labels = km.fit_predict(x)
```

labels

km.inertia_

```
plt.title('Data Distribution')
```

```
plt.xlabel('Annual Income')
```

```
plt.ylabel('Spending Score')
```

```
plt.scatter(x['Annual Income (k$)'], x['Spending Score (1-100)'], c=labels)
```

df[labels == 0]

df[labels == 0].max()

df[labels == 0].min()

ZERO = df[labels == 0]

ONE = df[labels == 1]

TWO = df[labels == 2]

THREE = df[labels == 3]

FOUR = df[labels == 4]

len(ZERO)

len(ONE)

len (TWO)

len (THREE)

len (FOUR)

len (FIVE)

len (SIX)

len (SEVEN)

len (EIGHT)

[ONE] width = 1

[TWO] width = 2

(y, 1) width = 1

(x, 2) width = 2

(x, 3) width = 3

Generally

len (x, y) = x + y + 1

Generally, when

(y, 1) width = 1

length is 1 then it is 1

9.

Linear Regression

ANSWER

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import statsmodels.api as sm

data = pd.read_csv ("D:\Subject\MLDLI-
                    - Datasets\Simple linear regression ol.csv")
data
data.describe()
# Our dependent variable (y) is the GPA
y = data ['GPA']
# Similarly, our independent variable (x)
x1 = data ['SAT']

plt.scatter (x1,y)
plt.xlabel ('SAT', fontsize = 20)
plt.ylabel ('GPA', fontsize = 20)
plt.show()

x = sm.add_constant (x1)
results = sm.OLS (y,x).fit()
results.summary()
plt.scatter (x1,y)
yhat = 0.0017 * x1 + 0.275
fig = plt.plot (x1,yhat, lw = 4, c = 'orange',
                label = 'regression line')
plt.xlabel ('SAT', fontsize = 20)

```

M	T	W	T	F	S	S
Page No.:	YOUVA					
Date:						

plt.ylabel('GPA', fontsize=20)
 plt.show()

10. Logistic Regression

Ans:-

```
import pandas as pd  
import numpy as np  
import statsmodels.api as sm  
import matplotlib.pyplot as plt  
import seaborn as sns  
sns.set()  
from scipy import stats  
stats.chisqprob = lambda chisq, df: stats.chi2.sf(chisq, df)
```

```
raw_data = pd.read_csv("D:\Subject"-  
- "MLDL\Datasets\Example-bank-data.csv")
```

raw_data

data = raw_data.copy()

data = data.drop(['Unnamed: 0'], axis=1)

data['y'] = data['y'].map({'yes': 1, -
- 'no': 0})

data

data.describe()

y = data['y']

x1 = data['duration']

x = sm.add_constant(x1)

reg_log = sm.Logit(y, x)

results_log = reg_log.fit()

results_log.summary()

plt.scatter(x1, y, color='co')

```

plt.xlabel('Duration', fontsize=20)
plt.ylabel('Subscription', fontsize=20)
plt.show()
    
```

```

np.set_printoptions(formatter={'float':lambda
    -da x: f'{x:0.2f}'}, format(x))
results_log.predict()
np.array(data['y'])
results_log.pred_table()
    
```

```

cm_df = pd.DataFrame(results_log.pred_table())
cm_df.columns = ['Predicted 0', 'Predicted 1']
cm_df = cm_df.rename(index={0: 'Actual 0',
                            -1: 'Actual 1'})
    
```

cm-df

```

cm = np.array(cm_df)
accuracy_train = (cm[0,0]+cm[1,1])/cm.sum()
    
```

accuracy_train

11. Multiple Linear Regression

Ans -

```
import numpy as np
import pandas as pd
import statsmodels.api as sm
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
```

```
raw_data = pd.read_csv("D:\Subject\ML-DL\Datasets\Dummy1.03+.csv")
raw_data['Attendance'] = raw_data['Attendance'].map({'Yes': 1, 'No': 0})
raw_data.describe()
```

~~y = raw_data['GPA']~~

~~x1 = raw_data[['SAT', 'Attendance']]~~

~~x = sm.add_constant(x1)~~

~~results = sm.OLS(y, x).fit()~~

~~results.summary()~~

~~plt.scatter(raw_data['SAT'], y)~~

~~yhat_no = 0.6439 + 0.0014 * raw_data['SAT']~~

~~yhat_yes = 0.2226 + 0.0014 * raw_data['SAT']~~

~~fig = plt.plot(raw_data['SAT'], yhat_no, -
lw=2, c='#006837')~~

~~fig = plt.plot(raw_data['SAT'], yhat_yes, -
lw=2, c='#a50026')~~

`plt.xlabel('SAT', fontsize=20)`
`plt.ylabel('GPA', fontsize=20)`
`plt.show()`

`plt.scatter(raw_data['SAT'], y, c=raw_data['-Attendance'], cmap='RdY1Gn-r')`

$$\hat{y}_{\text{no}} = 0.6439 + 0.0014 \cdot \text{raw_data}['SAT']$$

$$\hat{y}_{\text{yes}} = 0.2226 + 0.0014 \cdot \text{raw_data}['SAT']$$

`fig = plt.plot(raw_data['SAT'], yhat_no, lw=2, -c='#006837')`

`fig = plt.plot(raw_data['SAT'], yhat_yes, lw=2, -c='#a50026')`

`plt.xlabel('SAT', fontsize=20)`

`plt.ylabel('GPA', fontsize=20)`

`plt.show()`

`plt.scatter(raw_data['SAT'], raw_data['GPA'], c=raw_data['Attendance'], cmap='RdY1Gn-r')`

$$\hat{y}_{\text{no}} = 0.6439 + 0.0014 \cdot \text{raw_data}['SAT']$$

$$\hat{y}_{\text{yes}} = 0.2226 + 0.0014 \cdot \text{raw_data}['SAT']$$

$$\hat{y} = 0.0017 \cdot \text{raw_data}['SAT'] + 0.275$$

`fig = plt.plot(raw_data['SAT'], yhat_no, lw=2, -c='#006837', label='regression line1')`

`fig = plt.plot(raw_data['SAT'], yhat_yes, lw=2, -c='#a50026', label='regression line2')`

`fig = plt.plot(raw_data['SAT'], yhat, lw=3, -c='#4C72B0', label='regression line3')`

`plt.xlabel('SAT', fontsize=20)`

`plt.ylabel('GPA', fontsize=20)`

plt.show()

```
new_data = pd.DataFrame({ 'const': 1, 'SAT': -  
    - [1700, 1921], 'Attendance': [0, 1] })
```

```
new_data = new_data[[ 'const', 'SAT', 'Atten-  
    -dance' ]]
```

new_data

```
new_data.rename(index={0: 'Tom', 1: 'J-  
    -erry'})
```

```
predictions = results.predict(new_data)
```

predictions

```
predictionsdf = pd.DataFrame({ 'Predictions':  
    - : predictions })
```

```
joined = new_data.join(predictionsdf)
```

```
joined.rename(index={0: 'Tom', 1: 'Jerry'})
```

12. Predicting Customer churn using Decision Trees

Ans:-

```
import pandas as pd
from sklearn.model_selection import train_test-
    - split
```

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score,
    - classification_report
```

```
from sklearn import tree
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.datasets import make_classifi-
    - cation
```

```
X, y = make_classification(n_samples=1000, -
    - n_features=10, n_informative=5, n_redundant-
    - =2, n_classes=2, random_state=42)
```

```
df = pd.DataFrame(X, columns=[f'Feature_-
    - {i+1}' for i in range(X.shape[1])])
```

```
df['churn'] = y
```

```
print("Sample Data:")
```

```
print(df.head())
```

```
x_train, x_test, y_train, y_test = train_test-
    - split(df.drop(columns=['churn']), df-
    - ['churn'], test_size=0.3, random_state=42)
```

```
clf = DecisionTreeClassifier(max_depth=5, -
    - random_state=42)
```

```
clf.fit(x_train, y_train)
```

```
y-pred = clf.predict(X-test)
accuracy = accuracy_score(y-test, y-pred)
print ("In Accuracy: " + str(accuracy * 100) + "%")
print ("In Classification Report:")
print (classification_report(y-test, y-pred,
                             target_names = ["No Churn", "Churn"]))
plt.figure(figsize=(20,10))
tree.plot_tree(clf, feature_names=df.columns[-mms[-1]], class_names=["No Churn", "Churn"], filled=True)
plt.show()
```

13: Random forest model

Ans13

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

path = "https://archive.ics.uci.edu/ml/machine-
        learning-database/iris/iris.data"
headernames = ['sepal-length', 'sepal-width',
               'petal-length', 'petal-width', 'Class']
dataset = pd.read_csv(path, names=header-
                      -names)

dataset.head()
X = dataset.iloc[:, :-1].values
y = dataset.iloc[:, 4].values
X
y

from sklearn.model_selection import train-
        -test-split
X-train, X-test, y-train, y-test = train-test-
        -split(X, y, test_size=0.30)

from sklearn.ensemble import RandomForest-
        -Classifier
classifier = RandomForestClassifier(n_estimators=50)

classifier.fit(X-train, y-train)
RandomForestClassifier(n_estimators=50)
y-pred = classifier.predict(X-test)
```

```
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
result = confusion_matrix(y-test, y-pred)
print ("Confusion Matrix:")
print (result)
result1 = classification_report(y-test, y-pred)
print ("Classification_Report:")
print (result1)
result2 = accuracy_score(y-test, y-pred)
print ("Accuracy:", result2)
```

H.

SVM Demo1

Ans:-

```
import numpy as np  
import matplotlib.pyplot as plt  
import pandas as pd  
import matplotlib.pyplot as plt  
import sklearn
```

```
dataset = pd.read_csv('D:\\Subject\\MLDL-  
\\Datasets\\social-network-AdsI.csv')
```

```
dataset.head()
```

```
dataset.shape
```

```
x = dataset.iloc[:, [2, 3]]
```

```
y = dataset.iloc[:, 4]
```

```
x.head()
```

```
y.head()
```

```
from sklearn.model_selection import train-  
test_split
```

```
x_train, x_test, y_train, y_test = train-te-  
st-split(x, y, test_size=0.25, random_st-  
ate=0)
```

```
print("Training Data:", x_train.shape)
```

```
print("Testing Data:", x_test.shape)
```

```
from sklearn.preprocessing import Standard-  
Scaler
```

```
st_x = StandardScaler()
```

```
x_train = st_x.fit_transform(x_train)
```

```
x_test = st_x.transform(x_test)
```

from sklearn.svm import SVC
classifier = SVC(kernel='linear', random_state
= 0)

classifier.fit(x-train, y-train)

y-pred = classifier.predict(x-test)

y-pred

from sklearn.metrics import confusion_ma-
-trix

cm = confusion_matrix(y-test, y-pred)

from sklearn.metrics import accuracy_s-
-core

print("Accuracy: ", accuracy_score(-
[y-test, y-pred]))

cm

from sklearn.svm import SVC

classifier = SVC(kernel='rbf')

classifier.fit(x-train, y-train)

y-pred = classifier.predict(x-test)

y-pred

print("Accuracy: ", accuracy_score(-
[y-test, y-pred]))

from matplotlib.colors import ListedColormap

x-test x-set, y-set = x-train, y-train

x1, x2 = np.meshgrid(np.arange(start -
x-set[:, 0].min() - 1, stop = x-set[:, 0].-
max() + 1, step = 0.01),

```

nm.arange (start = x-set[:, 1].min() - 1, stop = -
(-x-set[:, 1].max() + 1, step = 0.01))

mtp.contourf (x1, x2, classifier.predict (nm.a-
-marray ([x1.ravel(), x2.ravel()]).T).reshape (-
-x1.shape), alpha = 0.75, cmap = ListedColormap (('red', -
('green'))))

mtp.xlim (x1.min(), x1.max())
mtp.ylim (x2.min(), x2.max())

for i, j in enumerate (nm.unique (y-set)):
    mtp.scatter (x-set[y-set == j, 0], x-set-
    [y-set == j, 1], c = ListedColormap (('red', 'green'))(i),
    label = j)

mtp.title ('SVM classifier (Training set)')
mtp.xlabel ('Age')
mtp.ylabel ('Estimated Salary')
mtp.legend ()
mtp.show()

```

```

from sklearn import matplotlib.colors
x-set, y-set = x-test, y-test
x1, x2, = nm.meshgrid (nm.arange (start =
- x-set[:, 0].min() - 1, stop = x-set[:, 0].max() + 1, step = 0.01),
nm.arange (start = x-set[:, 1].min() - 1,
- stop = x-set[:, 1].max() + 1, step = 0.01))
mtp.contourf (x1, x2, classifier.predict (n-
-m.array ([x1.ravel(), x2.ravel()]).T) - -

```

```
- reshape(x1.shape),  
alpha = 0.75, cmap = ListedColormap([{'red', -  
- 'green'}])  
mtp.xlim(x1.min(), x1.max())  
mtp.ylim(x2.min(), x2.max())  
for i, j in enumerate(np.unique(y_set)):  
    mtp.scatter(x_set[y_set == j, 0], x_set -  
    - [y_set == j, 1],  
    c = ListedColormap([{'red', 'green'}]) -  
    - (i), label = j)  
mtp.title('SVM classifier (Test set)')  
mtp.xlabel('Age')  
mtp.ylabel('Estimated Salary')  
mtp.legend()  
mtp.show()
```

15. SVM cancer Prediction

```
Ans) import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

Load Data from CSV file

```
cell-df = pd.read_csv("D:\Subject\MLDL-  
Datasets\cell-sample.csv")
```

```
cell-df.head()
```

```
cell-df.shape
```

```
cell-df.size
```

```
cell-df.count()
```

```
cell-df['class'].value_counts()
```

Distribution of the data set

```
benign-df = cell-df[cell-df['class']==2]-  
[0:200]
```

```
malignant-df = cell-df[cell-df['class']==4]-  
[0:200]
```

```
axes = benign-df.plot(kind='scatter', x=-  
'Clump', y='UnifSize', color='blue', label=  
'Benign')
```

```
malignant-df.plot(kind='scatter', x='Clu-  
mp', y='unifsize', color='red', label=  
'Benign').
```

```
axes = benign-df.plot(kind='scatter', x='Clump',  
y='unifsize', color='blue', label='Benign')
```

```
malignant-df.plot(kind='scatter', x='Clump',
```

- $y = \text{'Unifsize'}$, color = 'red', label = 'Benign', -
- ax=axes)

cell-df.dtypes

cell-df = cell-df[pd.to_numeric(cell-df=
- ['BareNue'], errors='coerce').notnull()]

cell-df['BareNue'] = cell-df['BareNue']-
- errors = 'coerce').notnull()

cell-df['BareNue'] = cell-df['BareNue']-
- errors = 'coerce').notnull()

cell-df['BareNue'] = cell-df['BareNue']-
- astype('int')

cell-df.dtypes

cell-df.columns

feature-df = cell-df[['Clump', 'Unifsize',
- 'MargAdh', 'SingEpisize', 'BareNue', -
- 'BludChrom', 'NormNucle', 'Mit']]

$x = \text{np.array(feature-df)}$

$y = \text{np.array(cell-df['class'])}$

$x[0:5]$

$y[0:5]$

from sklearn.model_selection import
- train-test-split

$x\text{-train}, x\text{-test}, y\text{-train}, y\text{-test} = \text{train-te-}$
- st-split(x, y, test_size=0.2, random=
- state=4)

$x\text{-train.shape}$

$y\text{-train.shape}$

$x\text{-test.shape}$

$y\text{-test.shape}$

from sklearn.svm import SVC

```

classifier = SVC(kernel='linear', random_state=0)
classifier.fit(x_train, y_train)

```

y-pred = classifier.predict(x-test)
y-pred

```

from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y-test, y-pred)
from sklearn.metrics import accuracy_score
print("Accuracy: ", accuracy_score(y-test, -y-pred))

```

```

from sklearn.svm import SVC
classifier = SVC(kernel='rbf')
classifier.fit(x_train, y_train)
y-pred = classifier.predict(x-test)
y-pred

```

```

print('Accuracy: ', accuracy_score(y-test, -y-pred))

```

```

from sklearn.metrics import classification_report
print(classification_report(y-test, y-pred))

```

The Sentiment Analysis using ANN

```
and import numpy as np  
import tensorflow as tf  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.preprocessing.text import Tokenizer  
from tensorflow.keras.layers import Dense, Embedding, GlobalAveragePooling1D  
from tensorflow.keras.preprocessing.sequence import pad_sequences  
from sklearn.model_selection import train_test_split
```

Sample Data: sentences with labels (1=positive, 0=negative)

sentences = [

"I love this product, it's amazing!",
"This is the worst thing I've ever bought.",
"Absolutely fantastic! Highly recommend.",
"I hate it, would not recommend to anyone.",
"It's okay, not the best but works fine.",
"I'm really satisfied with the quality.",
"Terrible experience, I'm very disappointed.",
"Great value for the money.",
"Not worth the price.",
"I'm so happy with this purchase!"

]

labels = np.array([1, 0, 1, 0, 1, 1, 0, 1, 0, 1])

Tokenize and Pad the sequence parameters

vocab-size = 50

embedding-dim = 16

max-length = 10

trunc-type = 'post'

padding-type = 'post'

oov-tok = ""

Tokenizer

tokenizer = Tokenizer(num_words=vocab-size,
- oov-token=oov-tok)

tokenizer.fit_on_texts(sentences)

word-index = tokenizer.word_index

sequences = tokenizer.texts_to_sequences(sentences)

padded-sequences = pad_sequences(sequences,
- maxlen=max-length, padding=padding-type,
- pre, truncating=trunc-type)

Train-Test Split

X-train, X-test, y-train, y-test = train-test-split-
- (padded-sequences, labels, test_size=0.2, rand-
- om-state=42)

Build the model

model = Sequential([

Embedding(vocab-size, embedding-dim, input-
- length=max-length),

GlobalAveragePooling1D(),

Dense(24, activation='relu'),

Dense(1, activation='sigmoid')

])

```
model.compile (loss='binary-crossentropy', -  
-optimizer='adam', metrics=[accuracy])  
model.summary()
```

epochs = 10

```
history = model.fit (X-train, y-train, epochs=  
- = epochs, validation_data=(X-test, y-test), -  
- verbose=2)
```

Evaluate the Model

```
loss, accuracy = model.evaluate (X-test, y-test,  
- verbose=2)
```

```
print ("Test Accuracy: {} accuracy * 100: -  
- .2f %.".format(accuracy))
```

Make Predictions

new-sentences = [

"I really love this!",

"I can't stand this at all."

]

new-sequences = tokenizer.texts_to_sequences(
- (new-sentences))

new-padded-sequences = pad_sequences (new-se-
-quences, maxlen=max-length, padding=pudd-
-ing-type, truncating=trunc-type)

Predict Sentiments

predictions = model.predict (new-padded-seq-
-uences)

for i, sentence in enumerate (new-sente-
-necs):

sentiment = "Positive" if predictions -

- [i] > 0.5 else "Negative"

~~print(f"sentences : {sentences} \n - {sentiment}")~~

~~notebook~~ 10/10/2021
16